




# Introduction to Clustering

# Clustering

- Classify objects (cases) into homogeneous groups called clusters.
- Objects in each cluster tend to be similar and dissimilar to objects in the other clusters.



# Computer vision application: Image segmentation





From: Image Segmentation by Nested Cuts, O. Veksler, CVPR2000


# What do we need for clustering?

## 1. Proximity measure, *either*

- similarity measure  $s(\mathbf{x}_i, \mathbf{x}_k)$ : large if  $\mathbf{x}_i, \mathbf{x}_k$  are similar
- dissimilarity(or distance) measure  $d(\mathbf{x}_i, \mathbf{x}_k)$ : small if  $\mathbf{x}_i, \mathbf{x}_k$  are similar



## 2. Criterion function to evaluate a clustering



*good clustering*



*bad clustering*

## 3. Algorithm to compute clustering


- For example, by optimizing the criterion function

# Distance (dissimilarity) measures

- Euclidean distance

$$d(\mathbf{x}_i, \mathbf{x}_j) = \sqrt{\sum_{k=1}^d (\mathbf{x}_i^{(k)} - \mathbf{x}_j^{(k)})^2}$$


- translation invariant



- Manhattan (city block) distance

$$d(\mathbf{x}_i, \mathbf{x}_j) = \sum_{k=1}^d |\mathbf{x}_i^{(k)} - \mathbf{x}_j^{(k)}|$$

- approximation to Euclidean distance, cheaper to compute



- They are special cases of **Minkowski distance**:

$$d_p(\mathbf{x}_i, \mathbf{x}_j) = \left( \sum_{k=1}^m |x_{ik} - x_{jk}|^p \right)^{\frac{1}{p}}$$

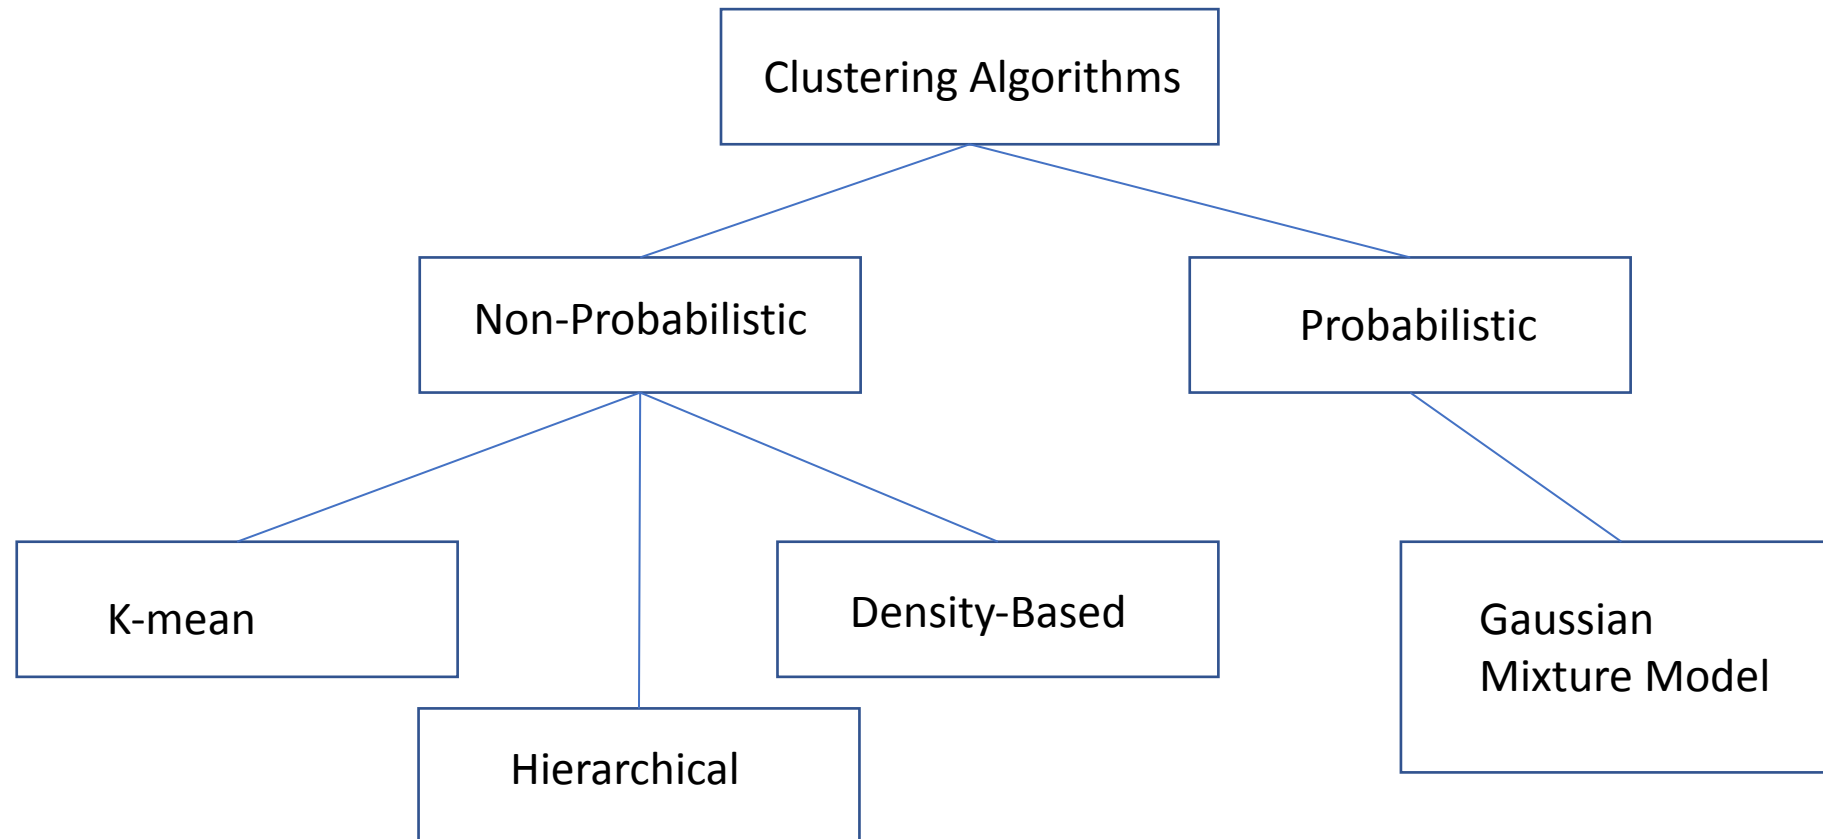
(p is a positive integer)

How about measures for mixed types?

# Clustering Algorithms

- Clustering Problem is an NP-Complete Problem
- Algorithms are Heuristic

# Clustering Algorithms (Cont.)





# K-Means Algorithm

- K-means (MacQueen, 1967) is a **partitional clustering** algorithm
- Let the set of data points  $D$  be  $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\}$ ,  
where  $\mathbf{x}_i = (x_{i1}, x_{i2}, \dots, x_{ir})$  is a **vector** in  $X \subseteq R^r$ , and  $r$  is the number of dimensions.
- The  $k$ -means algorithm partitions the given data into  $k$  clusters:
  - Each cluster has a cluster **center**, called **centroid**.
  - $k$  is specified by the user



# K-Means Algorithm (Cont.)

- Given  $k$ , the *k-means* algorithm works as follows:
  1. Choose  $k$  (random) data points (*seeds*) to be the initial *centroids*, cluster centers
  2. Assign each data point to the closest *centroid*
  3. Re-compute the *centroids* using the current cluster memberships
  4. If a convergence criterion is not met, repeat steps 2 and 3


# K-Means Convergence Criterion


- no (or minimum) re-assignments of data points to different clusters, *or*
- no (or minimum) change of centroids, *or*
- minimum decrease in the **sum of squared error** (SSE),


$$SSE = \sum_{j=1}^k \sum_{\mathbf{x} \in C_j} d(\mathbf{x}, \mathbf{m}_j)^2$$

- $C_j$  is the  $j$ th cluster,
- $\mathbf{m}_j$  is the centroid of cluster  $C_j$  (the mean vector of all the data points in  $C_j$ ),
- $d(\mathbf{x}, \mathbf{m}_j)$  is the (Euclidean) distance between data point  $\mathbf{x}$  and centroid  $\mathbf{m}_j$ .


# K-Means Example








Result of first iteration




Second iteration





Result of second iteration




# Why K-Means Algorithm?

- Strengths:
  - Simple: easy to understand and to implement
  - Efficient: Time complexity:  $O(tkn)$ ,  
where  $n$  is the number of data points,  
 $k$  is the number of clusters, and  
 $t$  is the number of iterations.
  - Since both  $k$  and  $t$  are small.  $k$ -means is considered a linear algorithm.
- K-means is the most popular clustering algorithm.
- Note that: it terminates at a **local optimum** if SSE is used. The **global optimum** is hard to find due to complexity.


# Weaknesses of K-Means Algorithm

- The algorithm is only applicable if the **mean** is defined.
  - For categorical data, *k*-mode - the centroid is represented by most frequent values.
- The user needs to specify *k*.
- The algorithm is sensitive to **outliers**
  - Outliers are data points that are very far away from other data points.
  - Outliers could be errors in the data recording or some special data points with very different values.

# Outliers



(A): Undesirable clusters




(B): Ideal clusters

# Outliers (Cont.)


- Remove some data points that are much further away from the centroids than other data points
  - To be safe, we may want to monitor these possible outliers over a few iterations and then decide to remove them.
- Perform random sampling: by choosing a small subset of the data points, the chance of selecting an outlier is much smaller
  - Assign the rest of the data points to the clusters by distance or similarity comparison, or classification




# Sensitivity to initial seeds




Random selection of seeds (centroids)




Random selection of seeds (centroids)




Iteration 1



Iteration 2




Iteration 1




Iteration 2

# Special data structures

- The  $k$ -means algorithm is not suitable for discovering clusters that are not hyper-ellipsoids (or hyper-spheres).



(A): Two natural clusters



(B):  $k$ -means clusters



# K-means summary

- Despite weaknesses,  $k$ -means is still the most popular algorithm due to its simplicity and efficiency
- No clear evidence that any other clustering algorithm performs better in general
- Comparing different clustering algorithms is a difficult task. No one knows the correct clusters!


# Hierarchical Clustering Methods

- **Hierarchical clustering** is characterized by the development of a hierarchy or tree-like structure.
  - **Agglomerative clustering** starts with each object in a separate cluster. Clusters are formed by grouping objects into bigger and bigger clusters.
  - **Divisive clustering** starts with all the objects grouped in a single cluster. Clusters are divided or split until each object is in a separate cluster.
- Agglomerative methods are commonly used in marketing research. They consist of linkage methods, variance methods, and centroid methods.

# Hierarchical Clustering Methods

- **Hierarchical clustering** is characterized by the development of a hierarchy or tree-like structure.
  - **Agglomerative clustering** starts with each object in a separate cluster. Clusters are formed by grouping objects into bigger and bigger clusters.
  - **Divisive clustering** starts with all the objects grouped in a single cluster. Clusters are divided or split until each object is in a separate cluster.
- Agglomerative methods are commonly used in marketing research. They consist of linkage methods, variance methods, and centroid methods.

# Dendrogram




# Agglomerative hierarchical clustering

initialize with each example in singleton cluster

**while** there is more than **1** cluster

1. find 2 nearest clusters
2. merge them



## ■ Four common ways to measure cluster distance

1. minimum distance  $d_{\min}(D_i, D_j) = \min_{x \in D_i, y \in D_j} \|x - y\|$

2. maximum distance  $d_{\max}(D_i, D_j) = \max_{x \in D_i, y \in D_j} \|x - y\|$

3. average distance  $d_{\text{avg}}(D_i, D_j) = \frac{1}{n_i n_j} \sum_{x \in D_i} \sum_{y \in D_j} \|x - y\|$


4. mean distance  $d_{\text{mean}}(D_i, D_j) = \|\mu_i - \mu_j\|$

# Hierarchical Agglomerative Clustering-Linkage Method

- The **single linkage** method is based on minimum distance, or the nearest neighbor rule.
- The **complete linkage** method is based on the maximum distance or the furthest neighbor approach.
- The **average linkage** method the distance between two clusters is defined as the average of the distances between all pairs of objects

# Single Link Example


---





# Complete Link Example

---



# Hierarchical Clustering Example

- Dataset -

Student_ID	Marks
1	10
2	7
3	28
4	20
5	35

# Hierarchical Clustering Example

## - Initial Step -

ID	1	2	3	4	5
1	0	3	18	10	25
2	3	0	21	13	28
3	18	21	0	8	7
4	10	13	8	0	15
5	25	28	7	15	0

Proximity Matrix



Initial Cluster

# Hierarchical Clustering Example

- 1<sup>st</sup> Merge Step -


ID	1	2	3	4	5
1	0	3	18	10	25
2	3	0	21	13	28
3	18	21	0	8	7
4	10	13	8	0	15
5	25	28	7	15	0



# Hierarchical Clustering Example


## - 2nd Merge Step -

ID	(1,2)	3	4	5
(1,2)	0	18	10	25
3	18	0	8	7
4	10	8	0	15
5	25	7	15	0




# Hierarchical Clustering Example

## - Final Clustering -



# Number of Clusters






# Density-based Clustering

- **Basic idea**

- Clusters are dense regions in the data space, separated by regions of lower object density
- A cluster is defined as a maximal set of density-connected points
- Discovers clusters of arbitrary shape

- **Method**

- DBSCAN




# Density Definition

- $\epsilon$ -Neighborhood – Objects within a radius of  $\epsilon$  from an object.

$$N_{\epsilon}(p) : \{q \mid d(p, q) \leq \epsilon\}$$

- “High density” -  $\epsilon$ -Neighborhood of an object contains at least *MinPts* of objects.




$\epsilon$ -Neighborhood of  $p$

$\epsilon$ -Neighborhood of  $q$

*Density of  $p$  is “high” (MinPts = 4)*

*Density of  $q$  is “low” (MinPts = 4)*

# Core, Border & Outlier



$\epsilon = 1\text{unit}$ ,  $\text{MinPts} = 5$


Given  $\epsilon$  and *MinPts*, categorize the objects into three exclusive groups.

A point is a **core point** if it has more than a specified number of points (MinPts) within Eps—These are points that are at the interior of a cluster.


A **border point** has fewer than MinPts within Eps, but is in the neighborhood of a core point.

A **noise point** is any point that is not a core point nor a border point.

# Example



Original Points




Point types: **core**,  
**border** and **outliers**

$\epsilon = 10$ , MinPts = 4



# Density-reachability

- Directly density-reachable
  - An object  $q$  is directly density-reachable from object  $p$  if  $p$  is a core object and  $q$  is in  $p$ 's  $\epsilon$ -neighborhood.




MinPts = 4

- $q$  is directly density-reachable from  $p$
- $p$  is not directly density-reachable from  $q$
- Density-reachability is asymmetric

# Density-reachability

- Density-Reachable (directly and indirectly):
  - A point  $p$  is directly density-reachable from  $p_2$
  - $p_2$  is directly density-reachable from  $p_1$
  - $p_1$  is directly density-reachable from  $q$
  - $p \leftarrow p_2 \leftarrow p_1 \leftarrow q$  form a chain




- $p$  is (indirectly) density-reachable from  $q$
- $q$  is not density-reachable from  $p$

# DBSCAN Algorithm: Example

- **Parameter**

- $\varepsilon = 2$  cm
- $MinPts = 3$




```
for each  $o \in D$  do
  if  $o$  is not yet classified then
    if  $o$  is a core-object then
      collect all objects density-reachable from  $o$ 
      and assign them to a new cluster.
    else
      assign  $o$  to NOISE
```

# DBSCAN Algorithm: Example

- Parameter

- $\varepsilon = 2$  cm
- $MinPts = 3$




```
for each  $o \in D$  do
  if  $o$  is not yet classified then
    if  $o$  is a core-object then
      collect all objects density-reachable from  $o$ 
      and assign them to a new cluster.
    else
      assign  $o$  to NOISE
```



# DBSCAN Algorithm: Example

- **Parameter**
  - $\varepsilon = 2$  cm
  - $MinPts = 3$



```
for each  $o \in D$  do
  if  $o$  is not yet classified then
    if  $o$  is a core-object then
      collect all objects density-reachable from  $o$ 
      and assign them to a new cluster.
    else
      assign  $o$  to NOISE
```

# DBSCAN: Sensitive to Parameters

Figure 8. DBScan results for DS1 with MinPts at 4 and Eps at (a) 0.5 and (b) 0.4.






Figure 9. DBScan results for DS2 with MinPts at 4 and Eps at (a) 5.0, (b) 3.5, and (c) 3.0.




## DBSCAN: Determining EPS and MinPts

- Idea is that for points in a cluster, their  $k^{\text{th}}$  nearest neighbors are at roughly the same distance
- Noise points have the  $k^{\text{th}}$  nearest neighbor at farther distance
- So, plot sorted distance of every point to its  $k^{\text{th}}$  nearest neighbor






## When DBSCAN Works Well




Original Points



Clusters


- Resistant to Noise
- Can handle clusters of different shapes and sizes

## When DBSCAN Does NOT Work Well




Original Points

- Cannot handle varying densities
- sensitive to parameters—hard to determine the correct set of parameters





(MinPts=4, Eps=9.92).



(MinPts=4, Eps=9.75)

# Internal Measures: SSE

- Clusters in more complicated figures aren't well separated
- Internal Index: Used to measure the goodness of a clustering structure without respect to external information
  - SSE
- SSE is good for comparing two clusterings or two clusters (average SSE).
- Can also be used to estimate the number of clusters



# Measure II: Cohesion and Separation

- **Cluster Cohesion**: Measures how closely related are objects in a cluster
  - Example: SSE
- **Cluster Separation**: Measure how distinct or well-separated a cluster is from other clusters
- Example: Squared Error
  - Cohesion is measured by the within cluster sum of squares (SSE)

$$WSS = \sum_i \sum_{x \in C_i} (x - m_i)^2$$

- Separation is measured by the between cluster sum of squares


$$BSS = \sum_i |C_i| (m - m_i)^2$$

- Where  $|C_i|$  is the size of cluster  $i$

# Internal Measures: Cohesion and Separation

- Example: SSE

- $BSS + WSS = \text{constant}$



**K=1 cluster:**

$$WSS = (1 - 3)^2 + (2 - 3)^2 + (4 - 3)^2 + (5 - 3)^2 = 10$$

$$BSS = 4 \times (3 - 3)^2 = 0$$

$$Total = 10 + 0 = 10$$

**K=2 clusters:**

$$WSS = (1 - 1.5)^2 + (2 - 1.5)^2 + (4 - 4.5)^2 + (5 - 4.5)^2 = 1$$

$$BSS = 2 \times (3 - 1.5)^2 + 2 \times (4.5 - 3)^2 = 9$$

$$Total = 1 + 9 = 10$$




# Measure III: Silhouette Coefficient

- Silhouette Coefficient combine ideas of both cohesion and separation, but for individual points, as well as clusters and clusterings
- For an individual point,  $i$ 
  - Calculate  $a(i)$  = average distance of  $i$  to the points in its cluster
  - Calculate  $b(i)$  = min (average distance of  $i$  to points in another cluster)
  - The silhouette coefficient for a point is then given by


$$s(i) = \frac{b(i) - a(i)}{\max\{a(i), b(i)\}}$$

- Typically between 0 and 1.
- The closer to 1 the better.




- Can calculate the Average Silhouette width for a cluster or a clustering


**Example:**




## K-Means (K=3)




## SSE vs Number of Clusters (Elbow Method)



## Silhouette Coefficient Plot (k=3)



## K-Means (K=2)



## Silhouette Coefficient Plot (k=2)

